

Test 4:

numplants = 167

targetsize = (300,300)

batchsize = 32

epochs = 100

Model: "sequential"

—		
Layer (type)	Output Shape	Param #
=====		
efficientnet-b3 (Model)	(None, None, None, 1536)	10783528
—		
global_averagepooling2d (Gl	(None, 1536)	0
—		
dense (Dense)	(None, 2048)	3147776
—		
dense1 (Dense)	(None, 1024)	2098176
—		
dense_2 (Dense)	(None, 167)	171175
=====		



Total params: 16,200,655

Trainable params: 16,113,359

Non-trainable params: 87,296

72/72 [=====] - 96s 1s/step - loss: 1.8357 - accuracy: 0.5551

test acc: 0.55512154



Conclusion: Accuracy was bad. Trying a deeper model. Going to test a different convolution base.